

Computer Vision Assignment 1

Omer Kamal Ali Ebead

Problem 1:

Preprocessing steps:


- Convert cifar-10 y labels to one hot vector so it can match the softmax output
- Normalising cifar-10 x images by dividing them by 255.0
- Try different architectures

1- Simple architecture:

```
Model: "sequential_9"
```

Layer (type)	Output Shape	Param #
conv2d_22 (Conv2D)	(None, 32, 32, 4)	112
average_pooling2d_20 (AveragePooling2D)	(None, 16, 16, 4)	0
conv2d_23 (Conv2D)	(None, 16, 16, 8)	520
average_pooling2d_21 (AveragePooling2D)	(None, 8, 8, 8)	0
conv2d_24 (Conv2D)	(None, 8, 8, 16)	2064
average_pooling2d_22 (AveragePooling2D)	(None, 4, 4, 16)	0
conv2d_25 (Conv2D)	(None, 4, 4, 32)	8224
average_pooling2d_23 (AveragePooling2D)	(None, 2, 2, 32)	0
flatten_9 (Flatten)	(None, 128)	0
dense_18 (Dense)	(None, 256)	33024
dense_19 (Dense)	(None, 128)	32896
dense_20 (Dense)	(None, 64)	8256
dense_21 (Dense)	(None, 10)	650

```
=====  
Total params: 85746 (334.95 KB)  
Trainable params: 85746 (334.95 KB)  
Non-trainable params: 0 (0.00 Byte)
```




- Around 55% validation accuracy.
- I needed to add more regularisation to the model and maybe make it more complex so the accuracy went higher.

2- Adding dropout and increase the number of layers:

Model: "sequential_10"

Layer (type)	Output Shape	Param #
conv2d_26 (Conv2D)	(None, 32, 32, 8)	608
max_pooling2d (MaxPooling2D)	(None, 16, 16, 8)	0
conv2d_27 (Conv2D)	(None, 16, 16, 16)	2064
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 16)	0
conv2d_28 (Conv2D)	(None, 8, 8, 32)	8224
dropout (Dropout)	(None, 8, 8, 32)	0
conv2d_29 (Conv2D)	(None, 8, 8, 64)	32832
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
conv2d_30 (Conv2D)	(None, 4, 4, 128)	131200
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 128)	0
dropout_1 (Dropout)	(None, 2, 2, 128)	0
conv2d_31 (Conv2D)	(None, 2, 2, 256)	524544
max_pooling2d_4 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten_10 (Flatten)	(None, 256)	0
dense_22 (Dense)	(None, 5012)	1288084
dense_23 (Dense)	(None, 256)	1283328
dense_24 (Dense)	(None, 10)	2570

Total params: 3273454 (12.49 MB)
 Trainable params: 3273454 (12.49 MB)
 Non-trainable params: 0 (0.00 Byte)




- Around 65% validation accuracy.

- Training and validation losses were decreasing until a point where it seems that the model is overfitting (7.5 epoch). I try early stopping next and Batch Normalisation.

3- Add Early stopping and Batch Normalisation:

Increasing the number of epochs to 100 and adding patience of 4 to the early stopping callback.

```
model3.summary()
Model: "sequential"
Layer (type)                Output Shape                Param #
-----
conv2d (Conv2D)              (None, 32, 32, 8)          392
conv2d_1 (Conv2D)            (None, 32, 32, 16)         2064
conv2d_2 (Conv2D)            (None, 32, 32, 32)         8224
batch_normalization (Batch Normalization) (None, 32, 32, 32)         128
dropout (Dropout)            (None, 32, 32, 32)         0
max_pooling2d (MaxPooling2D) (None, 16, 16, 32)         0
conv2d_3 (Conv2D)            (None, 16, 16, 64)         32832
conv2d_4 (Conv2D)            (None, 16, 16, 128)        131200
dropout_1 (Dropout)          (None, 16, 16, 128)        0
conv2d_5 (Conv2D)            (None, 16, 16, 256)        524544
batch_normalization_1 (Batch Normalization) (None, 16, 16, 256)        1024
max_pooling2d_1 (MaxPooling2D) (None, 8, 8, 256)         0
flatten (Flatten)            (None, 16384)              0
dense (Dense)                (None, 256)                4194560
dense_1 (Dense)              (None, 128)                32896
dense_2 (Dense)              (None, 64)                 8256
dense_3 (Dense)              (None, 10)                 650
=====
Total params: 4936770 (18.83 MB)
Trainable params: 4936194 (18.83 MB)
Non-trainable params: 576 (2.25 KB)
```



- Around 70% validation accuracy.
- The model stopped after the 10th epoch to prevent overfitting on the training data.

4- Add Snapshot Ensampling:

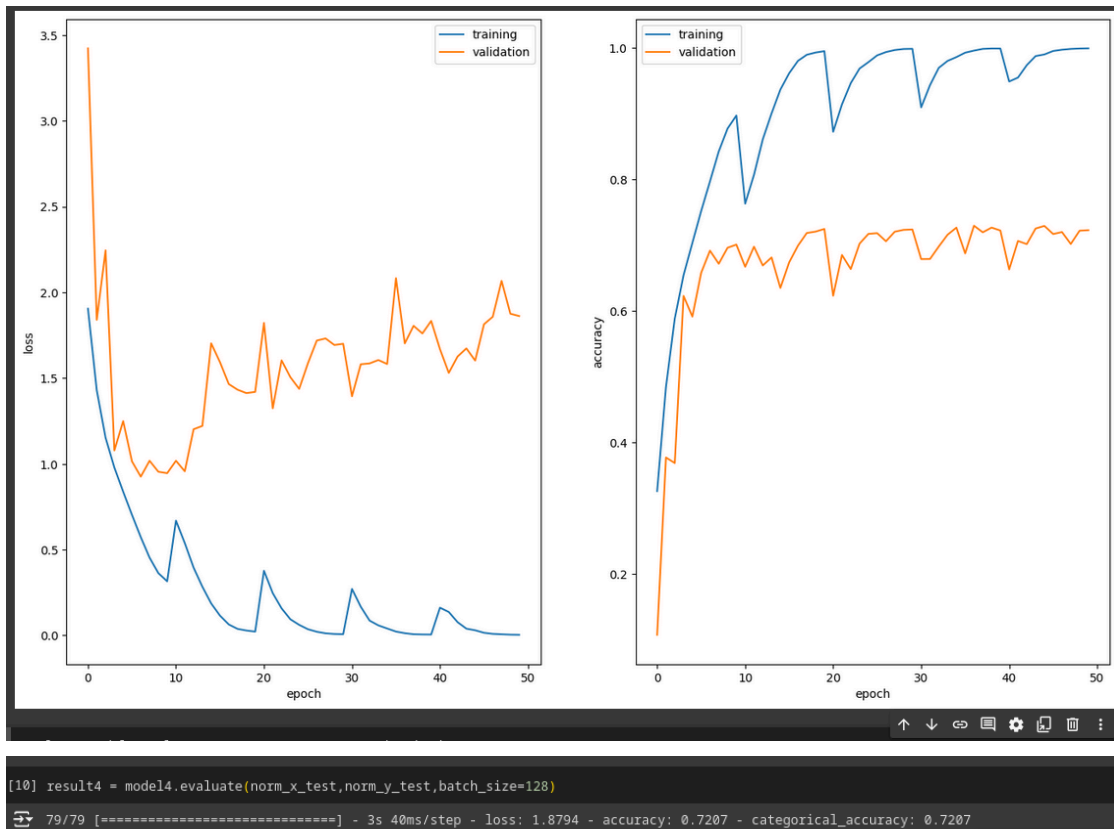
Paper: <https://arxiv.org/pdf/1704.00109v1>

Tutorial : <https://www.kaggle.com/code/fkdplc/snapshot-ensemble-tutorial-with-keras>

- I trained 5 models for 50 epochs (1 model per 10 epochs) and saved snapshots for each model (model parameters).

Model: "sequential"


Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 8)	392
dropout (Dropout)	(None, 32, 32, 8)	0
conv2d_1 (Conv2D)	(None, 32, 32, 16)	2064
dropout_1 (Dropout)	(None, 32, 32, 16)	0
conv2d_2 (Conv2D)	(None, 32, 32, 32)	8224
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	32832
batch_normalization_1 (Batch Normalization)	(None, 16, 16, 64)	256
dropout_2 (Dropout)	(None, 16, 16, 64)	0
conv2d_4 (Conv2D)	(None, 16, 16, 128)	131200
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 128)	512
dropout_3 (Dropout)	(None, 16, 16, 128)	0
conv2d_5 (Conv2D)	(None, 16, 16, 256)	524544
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 256)	1024
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 256)	0
flatten (Flatten)	(None, 16384)	0
dense (Dense)	(None, 256)	4194560
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 10)	650



- Validation accuracy differs from each model to the other but on average 70% accuracy.
- On testset it evaluates to 72%
- Maybe regularisation and increasing the number of epochs per model for can help optimise the results but I got caught by time.

Problem 2:

- I focused on the main three factors, operations, total accuracy and number of parameters. Therefore, I chose inception V3.
- I also considered choosing Inception V4 but I was not convinced that to achieve just a little higher accuracy, we need more operations and more parameters.




- Inception V3 received 256*256 pixels so i had to resize my dataset (32*32 pixels) using bilinear interpolation (take the average of the closest 4 pixels)
- I removed the Inception V3 model classifier and froze its feature extraction weights and added my classifier on top.
- The resulting architecture became :

Model: "model_1"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 256, 256, 3)]	0
inception_v3 (Functional)	(None, 6, 6, 2048)	21802784
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 2048)	0
dense_4 (Dense)	(None, 256)	524544
dense_5 (Dense)	(None, 128)	32896
dense_6 (Dense)	(None, 64)	8256
dense_7 (Dense)	(None, 10)	650

=====

Total params: 22369130 (85.33 MB)
Trainable params: 566346 (2.16 MB)
Non-trainable params: 21802784 (83.17 MB)



```

norm_x_test2 = resizing(norm_x_test)
result5 = model5.evaluate(norm_x_test2, norm_y_test, batch_size=128)



/79 [=====] - 81s 1s/step - loss: 0.5073 - accuracy: 0.8354 - categorical_accuracy: 0.8354

```


- Around 83% validation accuracy.
- On testset it evaluates to 83% also which is 10% better than my best model.

Problem 3:

- I made a function which receives the test image along with the dimensions of the occluder and slides it over the image.
- The function returns a list of images which are predicted by a pre-trained ResNet V2 model.
- I resized the test images so they can fit the model input layer with the same concept as in problem 2.
- Then I looped through the images and got predictions for each class.
- After that, I got only the probabilities of the correct class (toy terrier) from the predictions list.
- Below is a sample of the occluder sliding through the image with the prediction probability of the correct class "toy terrier".



- It is obvious that when we occlude the dog's face, the probability for the correct class reduces significantly.
- I created the below map to show the positions of the occluder with the probability of the model to classify the image correctly at that position.



Colab notebook: [🔗 omer_CV_Assignment_1.ipynb](#)